

---

## Context

The dataset `insurance.csv` contains anonymized information about a sample of individuals and their medical charges. Each entry includes demographic information, lifestyle factors, and the resulting medical costs for the individual. Key factors such as age, BMI, and smoking status are included, which will allow you to investigate patterns and correlations that impact healthcare expenses.

## Data dictionary

Please access the file `insurance.csv` from Canvas.

Column	Data type	Description
age	integer	Age of the primary beneficiary
gender	categorical	Gender of the beneficiary
bmi	categorical	Body Mass Index of the beneficiary
children	categorical	# of children covered by the insurance policy
smoker	binary	Indicates whether the beneficiary smoked or not
region	categorical	Residential area of the beneficiary in the US
charges	binary	Yearly medical costs billed by health insurance provider

Import the insurance data as a DataFrame and name it as `insurance`.

## Exploratory Data Analysis

1. How much more do smokers pay in medical charges compared to non-smokers across different age groups? To explore this, start by dividing the dataset into the following age groups: [18, 30), [30, 40), [40, 50), [50, 60), and [60, 70). For each age group, calculate the average medical charges separately for smokers and non-smokers. Create a bar chart that shows the average charges for smokers and non-smokers within each age group. Finally, identify the age group with the largest difference in charges between smokers and non-smokers, and report the value of this maximum difference.
2. Generate a correlation matrix for the numerical variables in the insurance dataset, including `age`, `bmi`, `children`, and `charges`, and visualize it using a heatmap. Identify which variable has the strongest correlation with `charges`.

---

## Linear Regression

3. Build a linear regression model using all available variables to predict charges (the target variable). Use `age`, `gender`, `bmi`, `children`, `smoker`, and `region` as predictor variables. After building the model, report which variables are statistically significant predictors of charges. Identify the significance level (e.g., 0.01, 0.05, or 0.1) for each significant variable.
4. Provide interpretations for the coefficients of `age`, `gender`, `bmi`, and `region`.

Note: The `region` variable has multiple (4) categories, resulting in three binary (dummy) variables in the model, each comparing a specific region to a reference region. Identify the reference region and interpret the coefficients of each dummy variable for region.

5. Let us now assess whether the model meets some of the core assumptions.
  - a. Create a histogram of residuals to visually inspect its distribution and comment on whether the residuals appear approximately normal.
  - b. Plot the residuals vs. the predicted charges to assess homoscedasticity and comment on whether the homoscedasticity assumption is met.
6. Create a new target variable named `log_charges` by applying a logarithmic transformation to the original `charges` variable. Refit the linear regression model using `log_charges` as the target variable instead of `charges`. For this model,
  - a. Interpret the coefficient of `age`.
  - b. Create a histogram of residuals to visually inspect its distribution and comment on whether the residuals appear approximately normal.
  - c. Plot the residuals vs. the predicted charges to assess homoscedasticity and comment on whether the homoscedasticity assumption is met.
7. Evaluate the overall fit of the model (developed in question 5) by reporting the R-squared and adjusted R-squared value. Explain the meaning of these statistics in the context of this model.
8. Using your model (developed in question 5), predict `charges` for five new observations based on the characteristics listed below:

<code>age</code>	<code>gender</code>	<code>bmi</code>	<code>children</code>	<code>smoker</code>	<code>region</code>
25	male	28.0	1	no	northeast
45	female	35.2	3	yes	southeast
32	male	30.5	0	no	northwest
54	female	24.7	2	yes	southwest
29	female	22.8	1	yes	southeast

---

## Logistic Regression

9. Let  $\mu$  denote the mean of `charges`. Create a new variable `binary_charges` as follows: For each observation, assign `binary_charges` a value of 0 if `charges` is less than or equal to  $\mu$  and a value of 1 if `charges` is greater than  $\mu$ . Report the % of observations assigned to each class (0 or 1) of `binary_charges`.
10. Build a logistic regression model using `age`, `gender`, `bmi`, `children`, `smoker`, and `region` as predictor variables and `binary_charges` as the target variable. After building the model, report which variables are statistically significant predictors of charges. Identify the significance level (e.g., 0.01, 0.05, or 0.1) for each significant variable.
11. Provide interpretations for the coefficients of `age`, `gender`, `bmi`, and `region`.

Note: The `region` variable has multiple (4) categories, resulting in three binary (dummy) variables in the model, each comparing a specific region to a reference region. Identify the reference region and interpret the coefficients of each dummy variable for region.

12. Using the logistic regression model predict `binary_charges` for five new observations based on the characteristics listed below:

age	gender	bmi	children	smoker	region
25	male	28.0	1	no	northeast
45	female	35.2	3	yes	southeast
32	male	30.5	0	no	northwest
54	female	24.7	2	yes	southwest
29	female	22.8	1	yes	southeast

## Decision Tree

13. Let  $q_1$  and  $q_3$  denote the first and third quartile of the charges variable. Create a new variable `multiclass_charges` as follows: For each observation, assign `multiclass_charges` a value of 'low' if `charges` is less than or equal to  $q_1$ , a value of 'medium' if `charges` is greater than  $q_1$  and less than or equal to  $q_3$ , and a value of 'high' if `charges` is greater than  $q_3$ . Report the % of observations assigned to each class ('low', 'medium', 'high') of `multiclass_charges`.
14. Build a decision tree model to predict `multiclass_charges` using `age`, `gender`, `bmi`, `children`, `smoker`, and `region` as predictor variables. Set a maximum depth of 4 for the decision tree to avoid overfitting. Report the number of leaf nodes in the tree.

- 
15. Plot the decision tree. Pick one of the paths from the root node to a leaf node and explain how the decision tree makes the classification decision on that path.
16. Using the decision-tree model predict `multiclass_charges` for five new observations based on the characteristics listed below:

age	gender	bmi	children	smoker	region
25	male	28.0	1	no	northeast
45	female	35.2	3	yes	southeast
32	male	30.5	0	no	northwest
54	female	24.7	2	yes	southwest
29	female	22.8	1	yes	southeast

- 
17. Create a column binary\_charges (as you did in Question 9). Build a logistic regression model (as you did in Question 10) using age, gender, bmi, children, smoker, and region as predictor variables and binary\_charges as the target variable. For this logistic regression model:
- Set a threshold (cut-off probability) of 0.2 and use the model to classify all the training observations. Report the confusion matrix.
  - Set a threshold of 0.8 and use the model to classify all the training observations. Report the confusion matrix.
  - Create an ROC plot for the logistic regression model and report the area under the curve (AUC). Explain the meaning of AUC in 1-2 lines.
18. Suppose we decide to use a threshold of 0.5 for the logistic regression model. Do the following:
- Report the confusion matrix.
  - Report the accuracy of the model and explain it (1-2 lines).
  - Report the precision of the model and explain it (1-2 lines).
  - Report the sensitivity of the model and explain it (1-2 lines).
  - Report the specificity of the model and explain it (1-2 lines).
  - Report the true positive rate of the model and explain it (1-2 lines).
  - Report the false positive rate of the model and explain it (1-2 lines).